

C_T

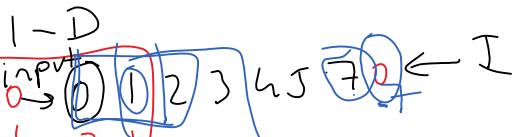
EE 417 Introduction to Computer Vision

/ EE 569 3D vision

Smoothing, Sharpening and Edges
Instructor: Associate Prof. Mehmet Keskinöz

Communication Theory & Technologies (CTT) Group,
Electronics Engineering Program,
Computer Science and Engineering,
Cyber-Security Program
Faculty of Engineering and Natural Sciences

Email: keskinoz@sabanciuniv.edu

Sabancı
Universitesi**C_T**

$$\max_k \rightarrow [1 \ 2 \ 1] \leftarrow M$$

$$I_{\text{new}} = \frac{I * M}{N} = \frac{4}{3}$$

C_T

Complexity of Convolution

$$\begin{array}{c} \xrightarrow{x} \\ \text{1-D} \\ I \leftarrow \text{input} \\ \text{Mask} \leftarrow G \\ \text{Input: } N \times N \\ \text{Mask: } M \times M \\ I_{\text{new}} = I * G = \sum_i G_i I(x+i) \end{array}$$

C_T

For 1-D
for each pixel (for each x)
we need M multiplication
,, , $M-1$ addition
Complexity $O(N \times M)$

C_T

$$\begin{array}{c} \text{In 2-D} \\ I \quad N \times N \\ M \quad M \times M \\ I_{\text{new}} = I * G = \sum_i \sum_j G_{ij} I(x+i, y+j) \end{array}$$

C_T

For each pixel ^{we need}
(M^2 multiplication
 $M^2 - 1$ addition
Complexity $O(N^2 M^2)$)

G_T

if the mask is separable
 $i.e., G_{ij} = f_i h_j$

Ex: $\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 \end{bmatrix}$

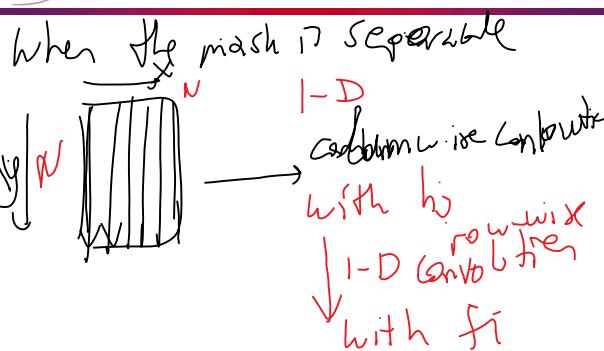
Row vector
 Column vector

G_T

if the mask is separable
 $I \cdot N \quad m \quad s$

$I_{new} = I * G = \sum_i \sum_j f_i(h_j) I(x_i, y_j)$

G_T



G_T

For each column $O(N^M)$
 \Rightarrow Column-wise convolution $O(N^2 M)$
 \Rightarrow Row-wise convolution $O(N^2 M)$
 \Rightarrow Total $O(2N^2 M)$

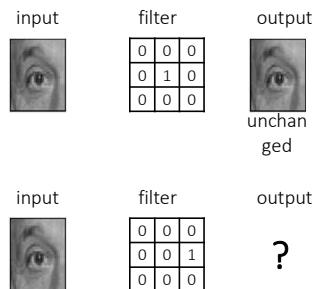
Other filters

input	filter ^{mask}	output
	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$?

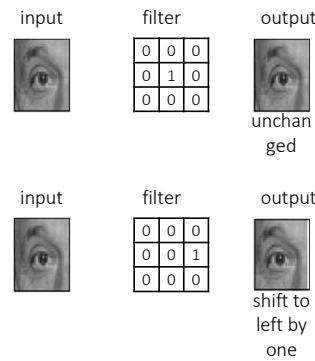
Other filters

input	filter	output
	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	 unchanged

Other filters

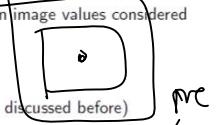


Other filters



Box Filter and Median Operator

Image smoothing aims at eliminating "outliers" in image values considered to be noise in a given context



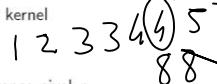
Box Filter (Mean Operator)

Local mean in a $(2k + 1) \times (2k + 1)$ window (as discussed before)

A simple option for image smoothing

Removes outliers but also reduces significantly the contrast

Often sufficient to use just a 3×3 or 5×5 filter kernel



Median Operator

Median of a $(2k + 1) \times (2k + 1)$ window to reference pixel p

Median of $2n + 1$ values = value at position $n + 1$ in sorted order

Example: 4 is median of $\{4, 7, 3, 1, 8, 7, 4, 5, 2, 3, 8\}$

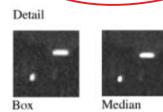
Removes outliers with only insignificant changes in image contrast



Example: Box Filter and Median Operator



Salt & Pepper noise



Detail



Bottom, left: 11×11 box filter. Bottom, right: 11×11 median operator



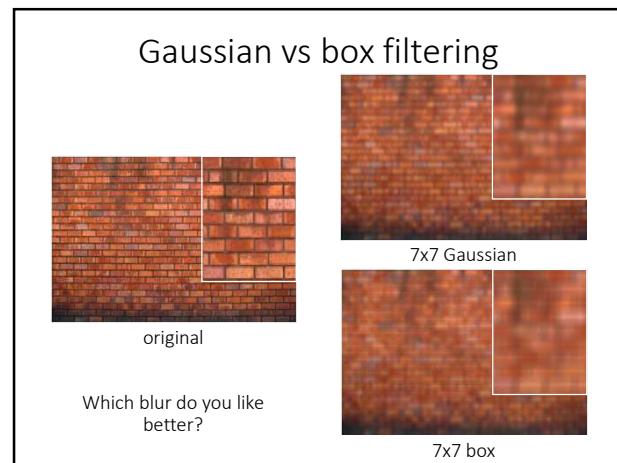
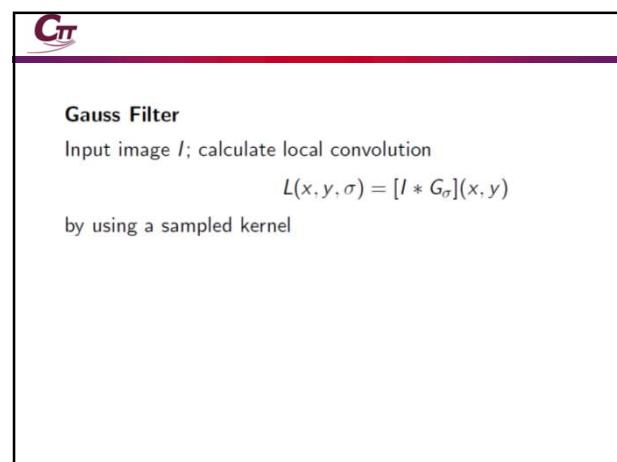
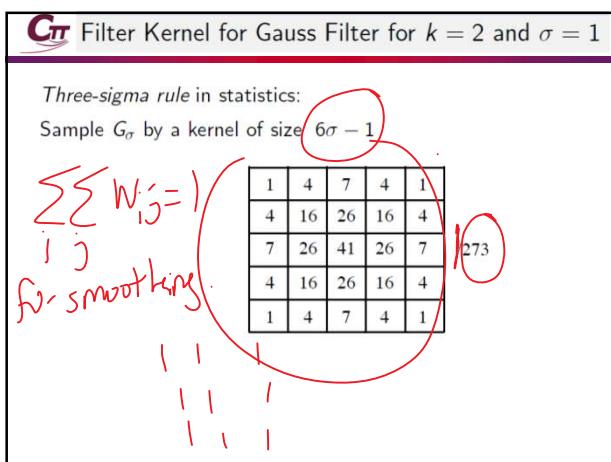
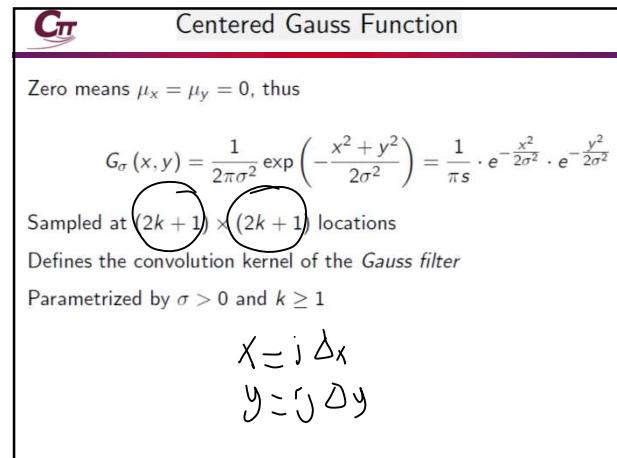
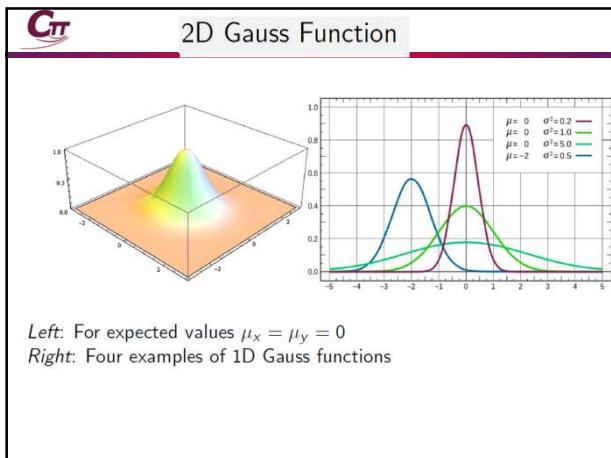
Gauss Filter



1-D Gaussian function

$$f(x) = e^{-\frac{(x - \mu)^2}{2\sigma^2}}$$

μ → center
 σ → standard deviation



Comparison of Noise Removal

Top, left: 128×128 input image with added uniform noise (± 15) Top, right: 3×3 box filter
 Bottom, left: 3×3 sigma-filter with $\sigma = 30$
 Bottom, middle: 3×3 median filter
 Bottom, right: Gauss filter for $\sigma = 1$ (as discussed above)

Sigma-Filter

Often useful local operator for noise removal

Assume a $(2k+1) \times (2k+1)$ window W_p

- ① Calculate the histogram of window W_p
- ② Calculate the mean μ of all values in the interval $[I(p) - \sigma, I(p) + \sigma]$
- ③ Let $J(p) = \mu$

Example: $\sigma = 50$ for $G_{\max} = 255$

Note:

$$\mu = \frac{1}{S} \cdot \sum_{u=l(p)-\sigma}^{l(p)+\sigma} u \cdot H(u)$$

$H(u)$ histogram value for W_p and $S = H(I(p) - \sigma) + \dots + H(I(p) + \sigma)$

*patch
= windowed image*

Sharpening

Sharpening aims at producing an enhanced image J by increasing the contrast of the given image I along edges, without adding too much noise within homogeneous regions in the image

I sharpens blurry

Unsharp Masking

- ① Residual $R(p) = I(p) - S(p)$ for smoothed version $S(p)$ of $I(p)$
- ② Add this residual to the given image I :

$$J(p) = I(p) + \lambda [I(p) - S(p)] = [1 + \lambda]I(p) - \lambda S(p)$$

where $\lambda > 0$ is a scaling factor

Any smoothing operator may be tried to produce $S(p)$

Parameter k controls the spatial distribution of the smoothing effect

Parameter λ controls the influence of the correction signal $[I(p) - S(p)]$

Examples of Sharpening

Illustration of unsharp masking for $k = 3$ and $\lambda = 1.5$

Upper left: 512×512 blurred input image

Upper right: Use of a median operator

Lower left: Use of a Gauss filter with $\sigma = 1$

Lower right: Use of a sigma filter with $\sigma = 25$

*sigma filter
T=25*

Image gradients

What are image edges?

$f(x)$

grayscale image

Very sharp discontinuities in intensity.

domain

$x = \begin{bmatrix} x \\ y \end{bmatrix}$

Detecting edges

How would you go about detecting edges in an image (i.e., discontinuities in a function)?

Detecting edges

How would you go about detecting edges in an image (i.e., discontinuities in a function)?

- ✓ You take derivatives: derivatives are large at discontinuities.

How do you differentiate a discrete image (or any other discrete signal)?

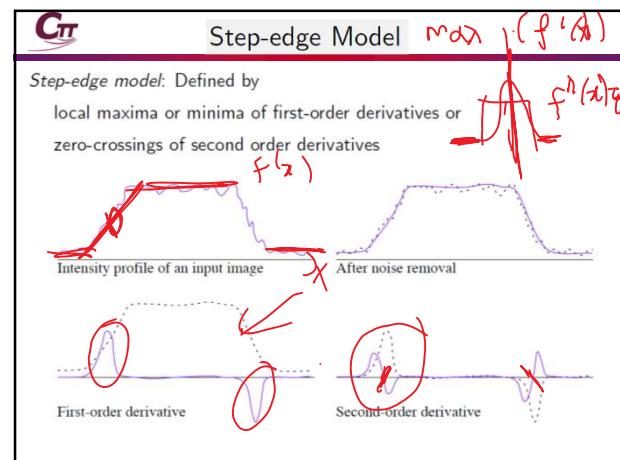
Detecting edges

How would you go about detecting edges in an image (i.e., discontinuities in a function)?

- ✓ You take derivatives: derivatives are large at discontinuities.

How do you differentiate a discrete image (or any other discrete signal)?

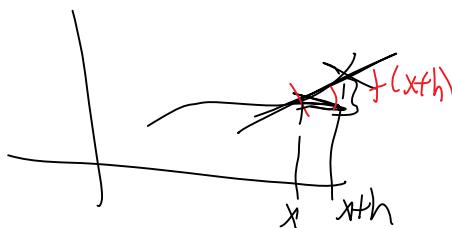
- ✓ You use finite differences.



Finite differences

High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$



Finite differences

High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

Alternative: use central difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+0.5h) - f(x-0.5h)}{h}$$

For discrete signals: Remove limit and set h=1

$$\boxed{0.5} \boxed{0} \boxed{0.5} f'(t) = \frac{f(t+1) - f(t-1)}{2}$$

What convolution kernel does this correspond to?

Finite differences

High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

Alternative: use central difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+0.5h) - f(x-0.5h)}{h}$$

For discrete signals: Remove limit and set h = 2

$$f'(x) = \frac{f(x+1) - f(x-1)}{2}$$

$$\begin{bmatrix} -1 & 0 & 1 \\ 1 & 0 & -1 \end{bmatrix}$$

Finite differences

High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

Alternative: use central difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+0.5h) - f(x-0.5h)}{h}$$

For discrete signals: Remove limit and set h = 2

$$f'(x) = \frac{f(x+1) - f(x-1)}{2}$$

1D derivative filter
 $\begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$

Cπ

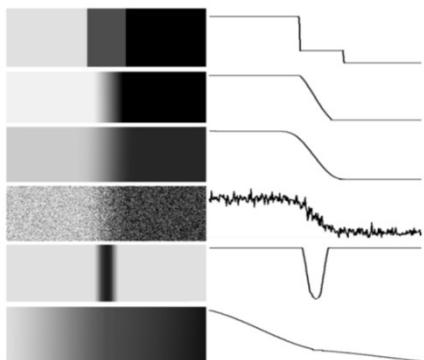
$$\vec{\nabla} f = \begin{bmatrix} f(x,y) \\ \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$$

Gradient

Cπ

$$\begin{aligned} \vec{\nabla} f(x,y) &= \begin{bmatrix} x^2y \\ 2xy \\ x^2 \end{bmatrix} \\ \frac{\partial f}{\partial x} &= 2xy \quad \vec{\nabla} f = \begin{bmatrix} x^2y \\ 2xy \\ x^2 \end{bmatrix} \\ \frac{\partial f}{\partial y} &= x^2 \end{aligned}$$

Image Profiles: Step Edges, a Line, and a Slope Change



Cπ

Derivatives

The derivative of a unary function f in the continuous case

$$\frac{df}{dx}(x) = f'(x) = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon) - f(x)}{\varepsilon}$$

Partial derivatives for function with two arguments, e.g.

$$\frac{\partial f}{\partial y}(x, y) = f_y(x, y) = \lim_{\varepsilon \rightarrow 0} \frac{f(x, y + \varepsilon) - f(x, y)}{\varepsilon}$$

But

In Ω we have the smallest distance $\varepsilon = 1$ between pixel locations
 No $\varepsilon \rightarrow 0$ possible in the image grid

Gr

Smoothing masks

$$\rightarrow \sum_i \sum_j g_{ij} = 1$$

Derivative mask

$$\sum_i \sum_j g_{ij} = D$$

Gr

Discrete Derivatives

Option 1: Reduce the derivative to a difference quotient for $\varepsilon = 1$

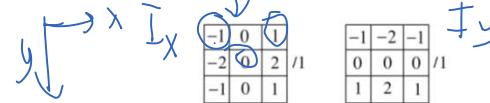
Option 2: Use a symmetric representation for $\varepsilon = 1$

$$I_y(x, y) = \frac{I(x, y+1) - I(x, y-1)}{2}$$

Thus no asymmetric bias, but very noise-sensitive

Option 3: Larger filter masks

For example, the Sobel operator uses the following kernels



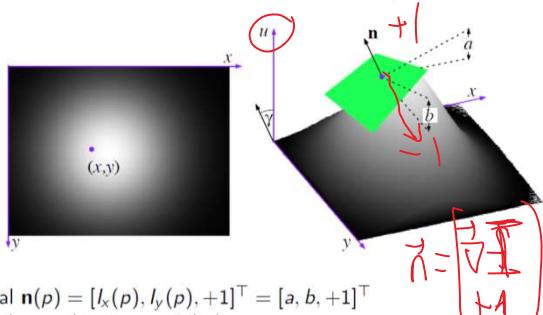
The diagram shows two 3x3 kernel matrices. The first matrix is labeled I_x and has values [-1, 0, 1; -2, 0, 2; -1, 0, 1] with a denominator of 1. The second matrix is labeled I_y and has values [-1, -2, -1; 0, 0, 0; 1, 2, 1] with a denominator of 1. Arrows point from the labels I_x and I_y to their respective matrices.

Gr

Derivatives Define Gradient and Normal

Gradient $[I_x(p), I_y(p)]^\top = [a, b]^\top$

Defined by partial derivatives $I_x(p)$ in x -, and $I_y(p)$ in y -direction



Normal $\mathbf{n}(p) = [I_x(p), I_y(p), +1]^\top = [a, b, +1]^\top$

Orthogonal to tangential plane

Gr

$\vec{\nabla} f \rightarrow$ magnitude
 \downarrow magnitude \rightarrow direction

$$\|\vec{\nabla} f\| = \sqrt{(\frac{\partial f}{\partial x})^2 + (\frac{\partial f}{\partial y})^2}$$

$$\vec{\nabla} f = +\hat{n}^{-1} \left(\frac{\partial f}{\partial n}, \frac{\partial f}{\partial x} \right)$$

↑ direction

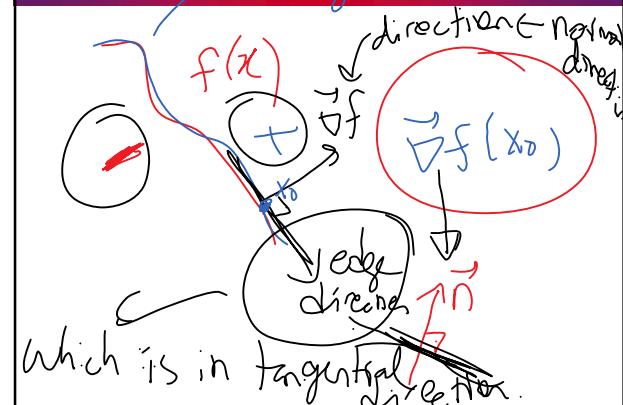
Gr

$$\nabla^- = [a, b]^\top$$

$$\nabla^+ = [b, -a]^\top$$

Gr

edge



The diagram illustrates the gradient flow and edge detection process. It shows a function $f(x)$ represented by a red curve. The gradient $\vec{\nabla} f$ is shown as a blue vector pointing upwards. The directient normal \vec{n} is shown as a red vector perpendicular to the gradient. A point x_0 is marked on the curve. The edge direction $\vec{\nabla} f(x_0)$ is shown as a blue vector pointing along the curve. A red circle highlights the edge direction. A red arrow points to the text "which is in tangential direction".

Gr

Gradient and Sobel Operator

Gradient vector $[I_x(p), I_y(p)]^\top$ at pixel location $p \in \Omega$

Magnitude of gradient

$$\sqrt{I_x(p)^2 + I_y(p)^2} = \|\text{grad } I(p)\|_2$$

(Note: $\|\cdot\|_2$ norm)

identifies local maxima or minima in $I_x(p)$ or $I_y(p)$ (i.e. edges)

Sobel Operator

- ① Convolution with the two masks shown on page before
- ② Produces at pixel p two values, $S_x(p)$ and $S_y(p)$
- ③ Sum $|S_x(p)| + |S_y(p)|$ is the result of the Sobel operator

This result approximates the magnitude of the gradient at p

Gr

$$\|\vec{x}\|_2 = \sqrt{\sum_i x_i^2}$$

Complex

$\Rightarrow L_2$ Norm

$$\|\vec{x}\|_1 = \sum_i |x_i|$$

Simple

$L_1 - \text{Norm}$

Gr

Example for Sobel Operator



Top: Original image I and absolute value of derivative in x -direction
Bottom: Absolute value of derivative in y -direction and final result
Results shown as inverted images (i.e. pixels with value 0 are white, grey pixels in between, and finally value 255 shown as black pixel).

Gr

Edge detection

- **Goal:** Identify sudden changes (discontinuities) in an image



- Intuitively, most semantic and shape information from the image can be encoded in the edges
- More compact than pixels

- **Ideal:** artist's line drawing (but artist is also using object-level knowledge)



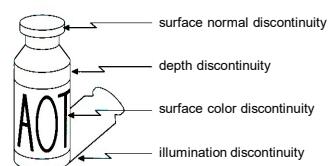
Source: D. Lowe and Debis

Gr

Edge detection



Origin of edges



Edges are caused by a variety of factors.

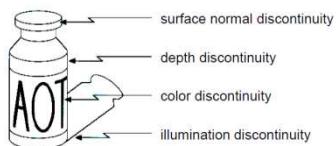
C₇

Geometric events

- surface orientation (boundary) discontinuities
- depth discontinuities
- color and texture discontinuities

Non-geometric events

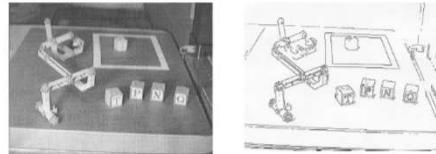
- illumination changes
- specularities
- shadows
- inter-reflections



C₇ Why is Edge Detection Useful?

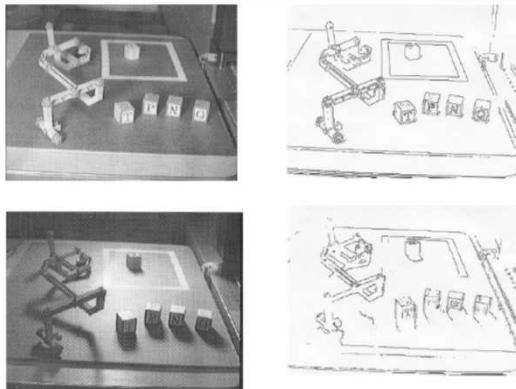
Important features can be extracted from the edges of an image (e.g., corners, lines, curves).

These features are used by higher-level computer vision algorithms (e.g., recognition).



C₇

Effect of Illumination

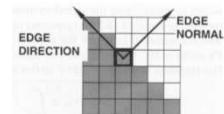


C₇

Edge Descriptors

Edge direction:

perpendicular to the direction of maximum intensity change (i.e., edge normal)



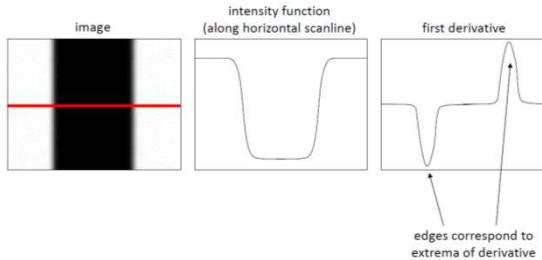
Edge strength: related to the local image contrast along the normal.

Edge position: the image position at which the edge is located.

C₇

Characterizing edges

- An edge is a place of rapid change in the image intensity function



C₇

Image gradient

The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$

$$\nabla f = [0, \frac{\partial f}{\partial y}]$$

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid increase in intensity

The gradient direction is given by

$$\theta = \tan^{-1} \left(\frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

The edge strength is given by the gradient magnitude

$$\| \nabla f \| = \sqrt{(\frac{\partial f}{\partial x})^2 + (\frac{\partial f}{\partial y})^2}$$

Source: Steve Seitz

Gr

Differentiation and convolution

Recall, for 2D function, $f(x,y)$:

$$\frac{\partial f}{\partial x} = \lim_{\epsilon \rightarrow 0} \left(\frac{f(x+\epsilon, y) - f(x, y)}{\epsilon} \right)$$

We could approximate this as:

$$\frac{\partial f}{\partial x} \approx \frac{f(x_{n+1}, y) - f(x_n, y)}{\Delta x}$$

(which is obviously a convolution)

-1	1
----	---

Check!

Source: D. Forsyth, D. Lowe

Computing image gradients

- Select your favorite derivative filters.

$$S_x = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$S_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

- Convolve with the image to compute derivatives.

$$\frac{\partial f}{\partial x} = S_x \otimes f \quad \frac{\partial f}{\partial y} = S_y \otimes f$$

Computing image gradients

- Select your favorite derivative filters.

$$S_x = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} \quad S_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

- Convolve with the image to compute derivatives.

$$\frac{\partial f}{\partial x} = S_x \otimes f \quad \frac{\partial f}{\partial y} = S_y \otimes f$$

- Form the image gradient, and compute its direction and amplitude.

$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x} & \frac{\partial f}{\partial y} \end{bmatrix} \quad \theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right) \quad \| \nabla f \| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

gradient

direction

amplitude

Gr

4.) Determine gradient magnitude

$L_1 : \| S_x \|_1 + \| S_y \|_1$

norm

$L_2 : \| S_x \|_2 + \| S_y \|_2$

Several derivative filters

Sobel

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

Scharr

$$\begin{bmatrix} 3 & 0 & -3 \\ 10 & 0 & -10 \\ 3 & 0 & -3 \end{bmatrix}$$

Roberts

$$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 1 & 0 \\ -1 & 0 & 1 \end{bmatrix}$$

- How are the other filters derived and how do they relate to the Sobel filter?
- How would you derive a derivative filter that is larger than 3x3?

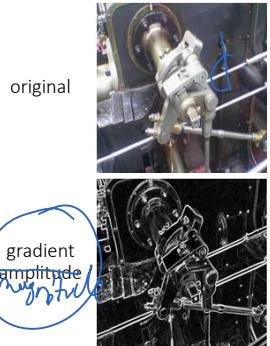
Computing image gradients

- Select your favorite derivative filters.

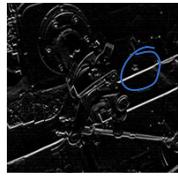
$$S_x = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$S_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Image gradient example



vertical derivative

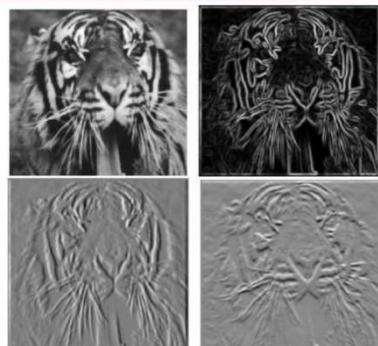


horizontal derivative



How does the gradient direction relate to these edges?

Finite differences: example



Which one is the gradient in the x-direction (resp. y-direction)?

The Sobel filter

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

Sobel filter

= * What filter is this?

$$\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix}$$

1D derivative filter

The Sobel filter

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

Sobel filter

= * Blurring

$$\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix}$$

1D derivative filter

In a 2D image, does this filter respond along horizontal or vertical lines?

The Sobel filter

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

Sobel filter

= * Blurring

$$\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix}$$

1D derivative filter

Does this filter return large responses on vertical or horizontal lines?

The Sobel filter

Horizontal Sobel filter:

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

= * $\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix}$

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

What does the vertical Sobel filter look like?

The Sobel filter

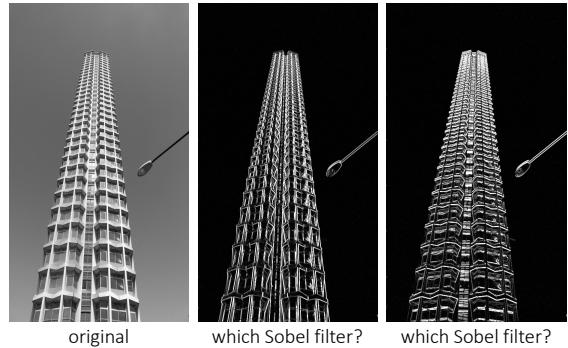
Horizontal Sobel filter:

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} = \begin{pmatrix} 1 \\ 2 \\ 1 \end{pmatrix} * \begin{pmatrix} 1 & 0 & -1 \end{pmatrix}$$

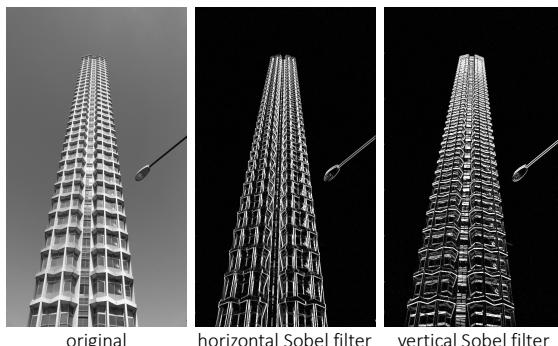
Vertical Sobel filter:

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} = \begin{pmatrix} 1 \\ 0 \\ -1 \end{pmatrix} * \begin{pmatrix} 1 & 2 & 1 \end{pmatrix}$$

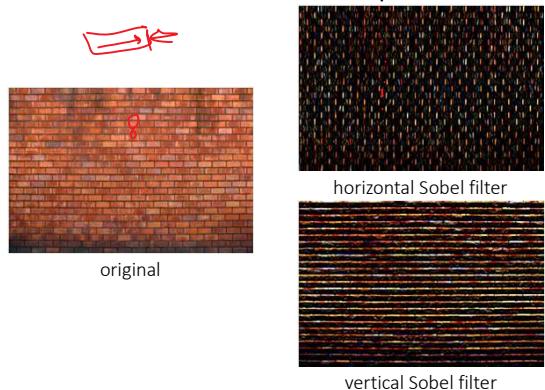
Sobel filter example



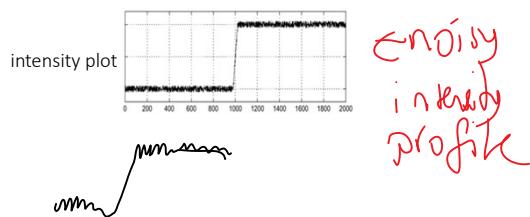
Sobel filter example



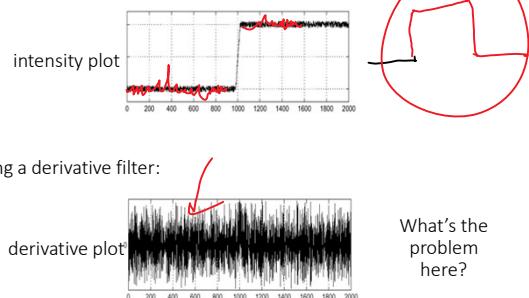
Sobel filter example



How do you find the edge of this signal?



How do you find the edge of this signal?



C_T

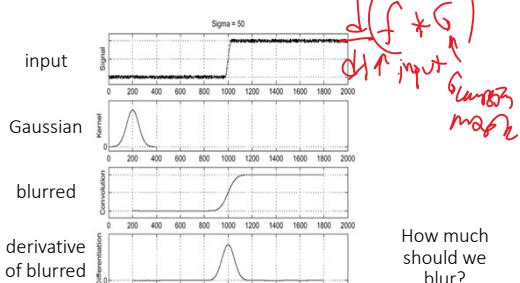
Effects of noise

- Finite difference filters respond strongly to noise
 - Image noise results in pixels that look very different from their neighbors
 - Generally, the larger the noise the stronger the response
- What is to be done?
 - Smoothing the image should help, by forcing pixels different to their neighbors (=noise pixels?) to look more like neighbors

Source: D. Forsyth

Differentiation is very sensitive to noise

When using derivative filters, it is critical to blur first!



How much should we blur?

C_T

Derivative theorem of convolution

- Differentiation and convolution both linear operators: they "commute"

$$\frac{d}{dx}(f * g) = \frac{df}{dx} * g = f * \frac{dg}{dx}$$

(Note: $\frac{dg}{dx}$ is circled in red)
- This saves us one operation:

Source: S. Seitz

C_T

md ximmo mminus

$\nabla f \leftarrow \text{gradient}$

$\frac{\partial f}{\partial x} \quad \frac{\partial^2 f}{\partial x^2}$

$\frac{\partial^2 f}{\partial x^2} = 0$

$\frac{\partial^2 f}{\partial x^2} \quad \text{Laplace}$

C_T

$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$ *vector, magnitude direction*

$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$ *scalar magnitude*

No directional information. Isotropic

Laplace filter

Basically a second derivative filter.

- We can use finite differences to derive it, as with first derivative filter.

first-order finite difference $f'(x) = \lim_{h \rightarrow 0} \frac{f(x+0.5h) - f(x-0.5h)}{h}$ \rightarrow 1D derivative filter $\begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$

second-order finite difference $f''(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - 2f(x) + f(x-h)}{h^2}$ \rightarrow Laplace filter $\begin{bmatrix} 1 & -2 & 1 \end{bmatrix}$?

Laplace filter

Basically a second derivative filter.

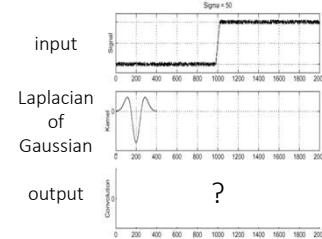
- We can use finite differences to derive it, as with first derivative filter.

$$\text{first-order finite difference } f'(x) = \lim_{h \rightarrow 0} \frac{f(x+0.5h) - f(x-0.5h)}{h} \rightarrow \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$

$$\text{second-order finite difference } f''(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - 2f(x) + f(x-h)}{h^2} \rightarrow \begin{bmatrix} 1 & -2 & 1 \end{bmatrix}$$

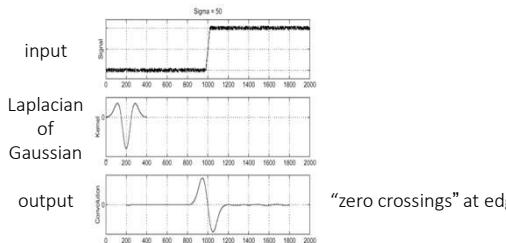
Laplacian of Gaussian (LoG) filter

As with derivative, we can combine Laplace filtering with Gaussian filtering

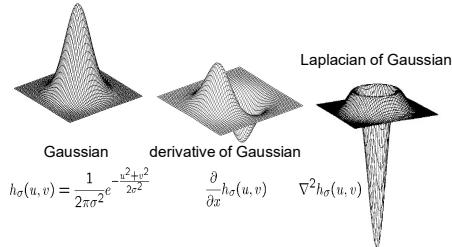


Laplacian of Gaussian (LoG) filter

As with derivative, we can combine Laplace filtering with Gaussian filtering



2D edge detection filters

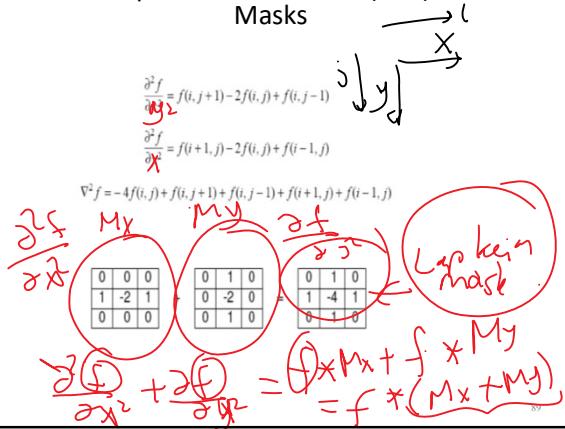


∇^2 is the Laplacian operator:

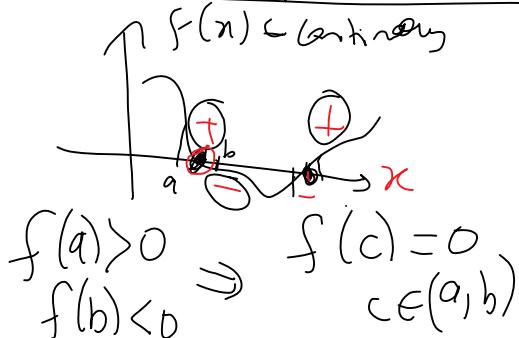
$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

88

Laplacian of Gaussian (LoG) Masks



zero crossing of a 1-D function



90

$$\downarrow \frac{\partial^2 f}{\partial x^2} \rightarrow \text{mask}$$

91

Laplacian of Gaussian (LoG) Masks

$$\nabla^2(f * g) = f * (\nabla^2 g) \rightarrow \text{LoG}$$

gruss: a

0	1	0	1
1	-4	1	1
0	1	0	1

1	1	1	1
1	-8	1	1
1	1	1	1

-1	2	-1	1
2	-4	2	1
-1	2	-1	1

92

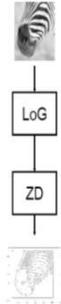
Using the LoG Function (Laplacian of Gaussian)

- The LoG function will be
 - Zero far away from the edge
 - Positive on one side
 - Negative on the other side
 - Zero just at the edge
- It has simple digital mask implementation(s)
- So it can be used as an edge operator

93

Edge detection strategy

- filter with Laplacian of Gaussian
- detect zero crossings
- mark the zero points where:
 - there is a sufficiently large derivative,
 - and enough contrast
- once again we have parameters
 - scale of Gaussian smoothing
 - thresholds
- once again no set of universal parameters
- does not seem to be better than the strategy of looking for maxima of gradient magnitude.

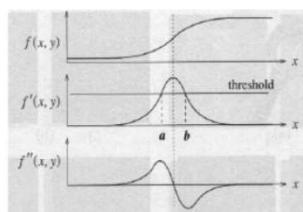


94

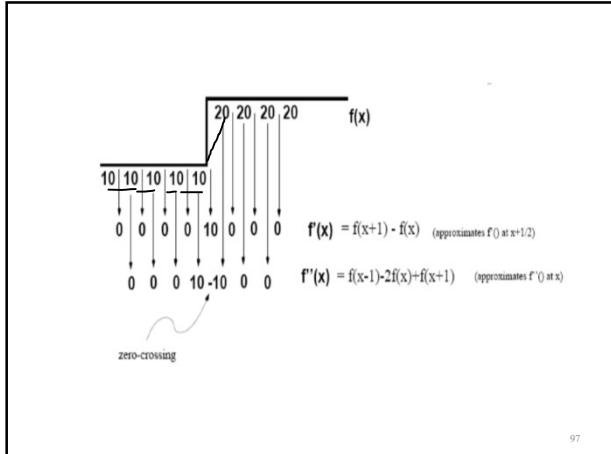
- Approximate finding maxima/minima of gradient magnitude by finding places where:
$$\frac{df^2}{dx^2}(x) = 0$$
- Can't always find discrete pixels where the second derivative is zero – look for zero-crossing instead.

95

- Edge points can be detected by finding the zero-crossings of the second derivative.



96



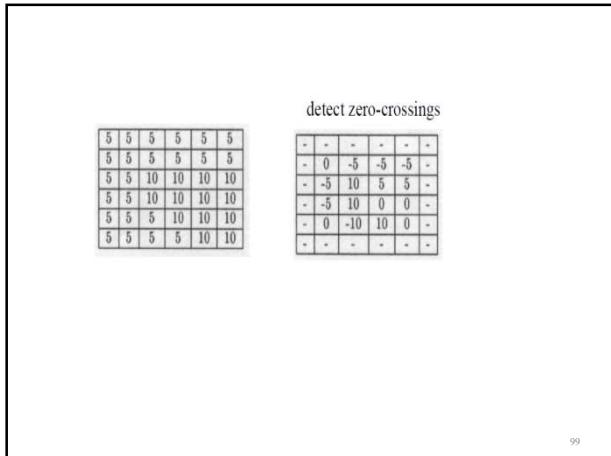
- Four cases of zero-crossings:



- Slope of zero-crossing $\{a, -b\}$ is: $|a+b|$.

- To detect “strong” zero-crossing, threshold the slope.

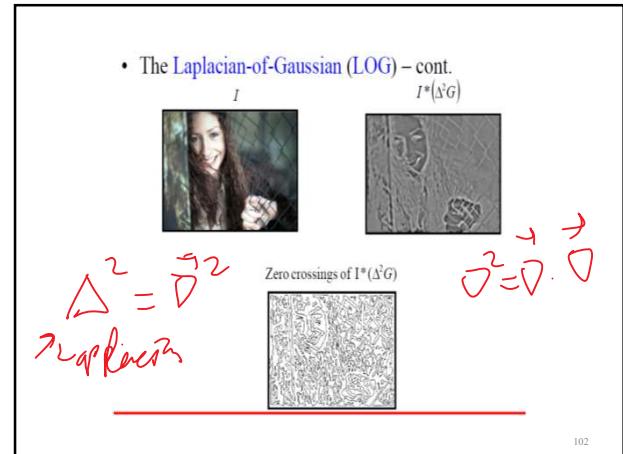
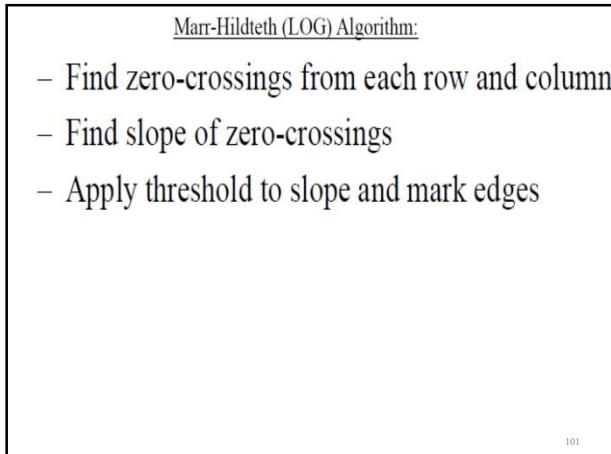
$|slope| > \text{threshold}$

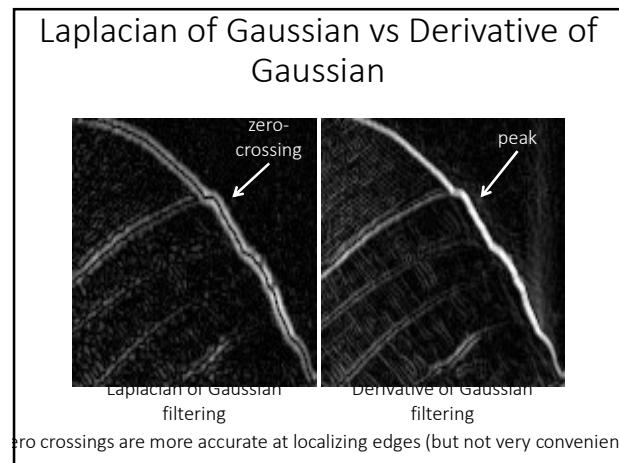
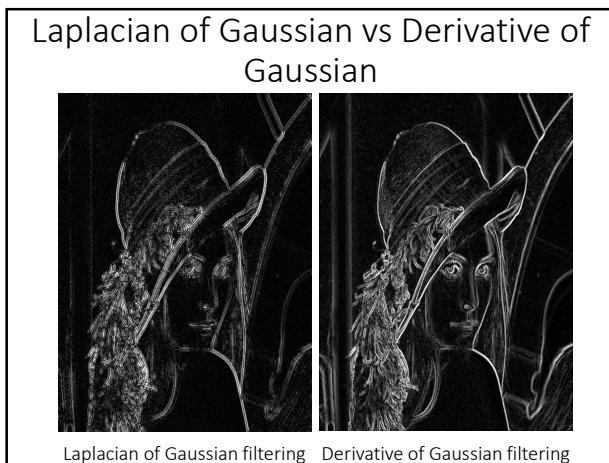
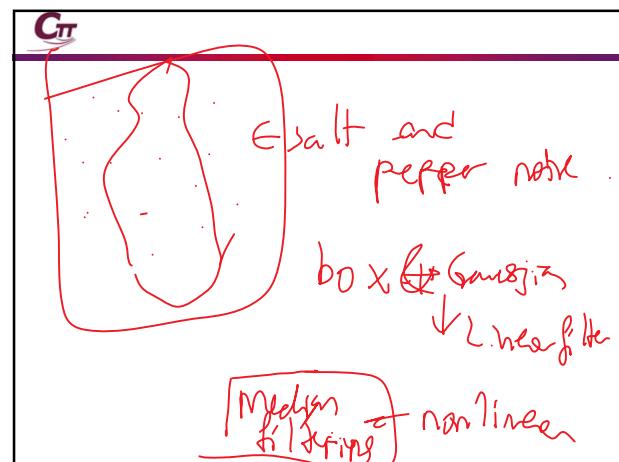
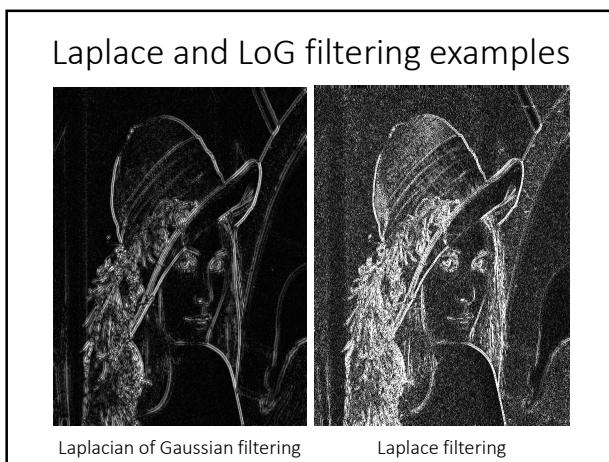
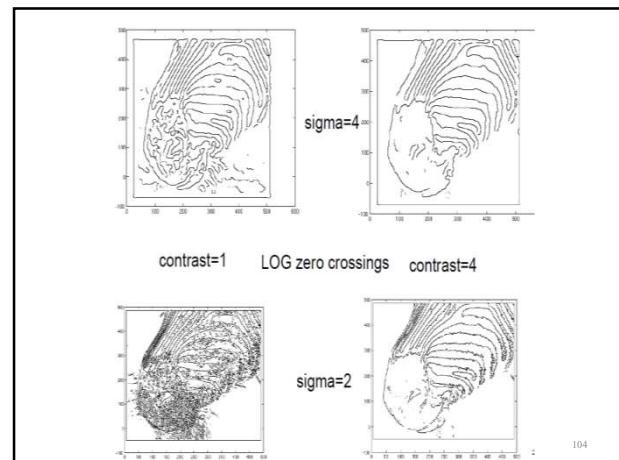
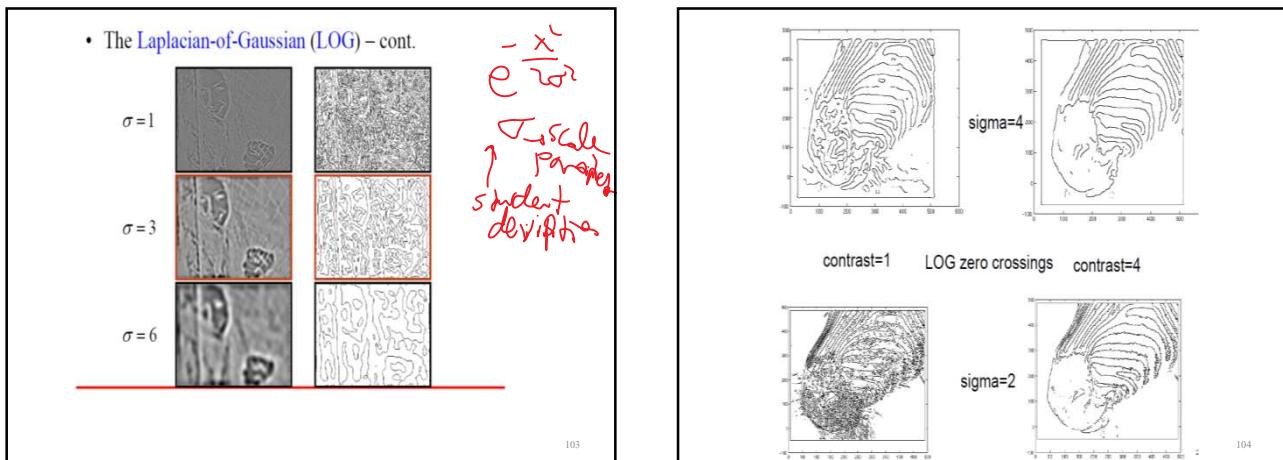


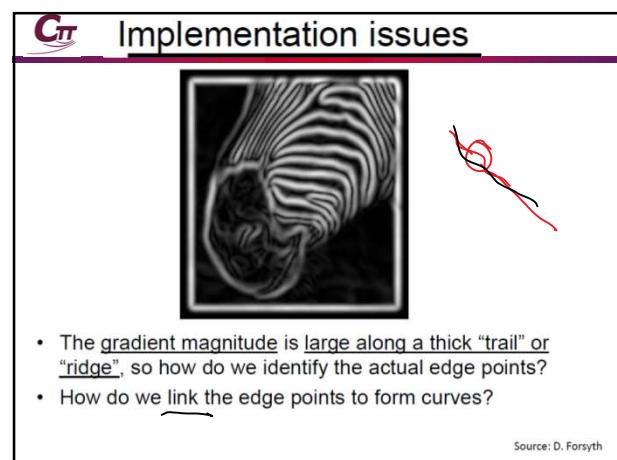
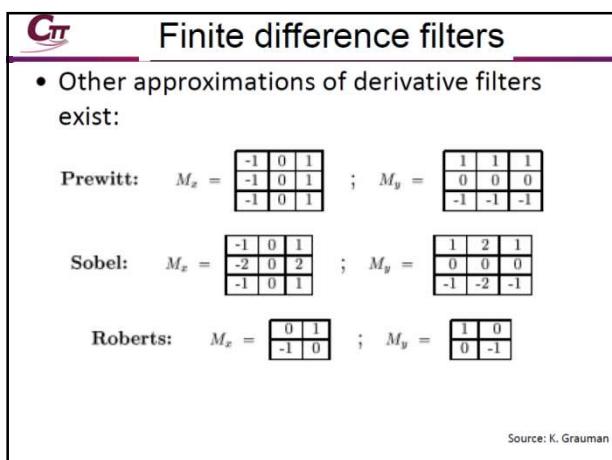
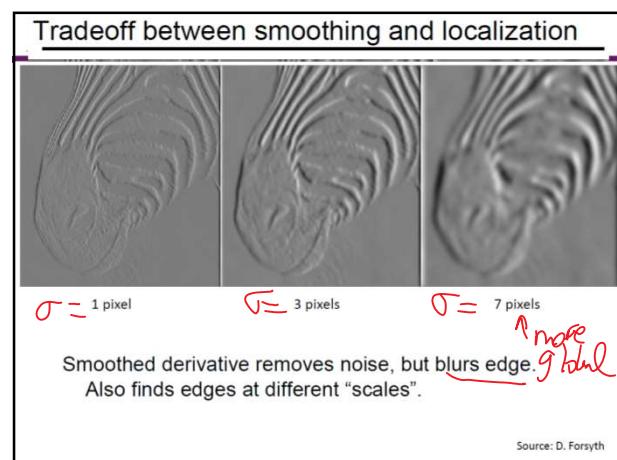
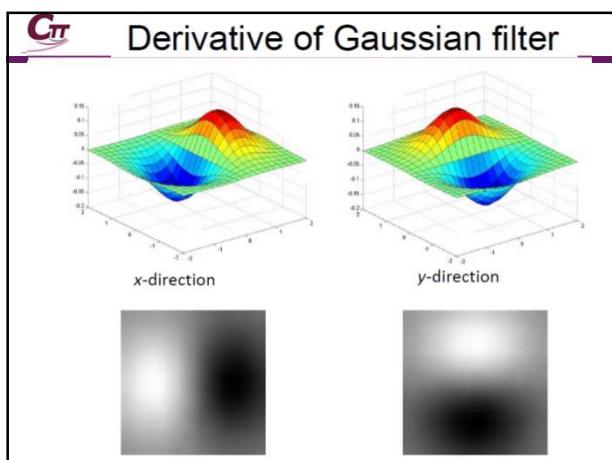
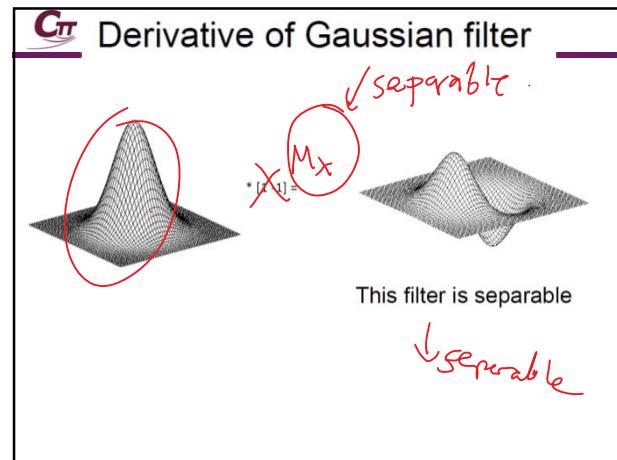
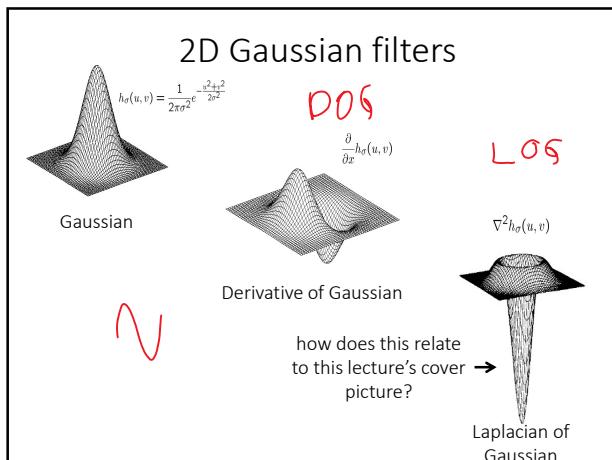
Properties of Laplacian

- It is cheaper to implement than the gradient (i.e., one mask only).
- It does not provide information about edge direction.
- It is more sensitive to noise (i.e., differentiates twice).
- It is an isotropic operator (equal weights in all directions)

100



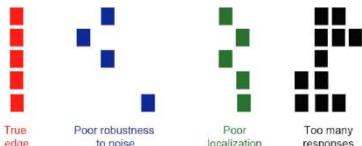




C_π

Designing an edge detector

- Criteria for an “optimal” edge detector:
 - Good detection:** the optimal detector must minimize the probability of **false positives** (detecting spurious edges caused by noise), as well as that of **false negatives** (missing real edges)
 - Good localization:** the edges detected must be as close as possible to the **true edges**
 - Single response:** the detector must return **one point** only for each true edge point; that is, minimize the number of local maxima around the true edge



Source: L. Fei-Fei

C_π

Canny edge detector

It is a multi-stage algorithm and we will go through each stages.

1. Noise Reduction

Since edge detection is susceptible to noise in the image, first step is to remove the noise in the image with a 5x5 Gaussian filter.

C_π

2. Finding Intensity Gradient of the Image

Smoothed image is then filtered with a Sobel kernel in both horizontal and vertical direction to get first derivative in horizontal direction (G_x) and vertical direction (G_y). From these two images, we can find edge gradient and direction for each pixel as follows:

$$\text{Edge Gradient } (G) = \sqrt{G_x^2 + G_y^2}$$

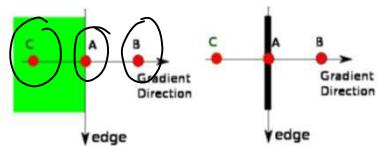
$$\text{Angle } (\theta) = \tan^{-1} \left(\frac{G_y}{G_x} \right)$$

Gradient direction is always perpendicular to edges. It is rounded to one of four angles representing vertical, horizontal and two diagonal directions.

C_π

3. Non-maximum Suppression

After getting gradient magnitude and direction, a full scan of image is done to remove any unwanted pixels which may not constitute the edge. For this, at every pixel, pixel is checked if it is a local maximum in its neighborhood in the direction of gradient. Check the image below:



Point A is on the edge (in vertical direction). Gradient direction is normal to the edge. Point B and C are in gradient directions. So point A is checked with point B and C to see if it forms a local maximum. If so, it is considered for next stage, otherwise, it is suppressed (put to zero).

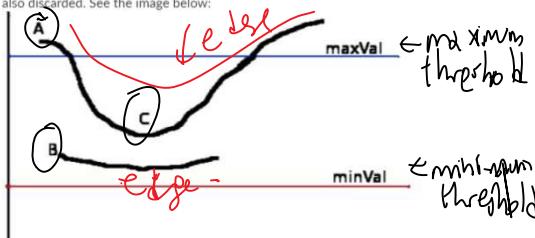
you get “thin edges”.

non-maximum suppression to get rid of spurious response to edge detection

C_π

4. Hysteresis Threshholding

This stage decides which are all edges are really edges and which are not. For this, we need two threshold values, minVal and maxVal . Any edges with intensity gradient more than maxVal are sure to be edges and those below minVal are sure to be non-edges, so discarded. Those who lie between these two thresholds are classified edges or non-edges based on their connectivity. If they are connected to “sure-edge” pixels, they are considered to be part of edges. Otherwise, they are also discarded. See the image below:



C_π

The edge A is above the maxVal , so considered as “sure-edge”. Although edge C is below maxVal , it is connected to edge A, so that also considered as valid edge and we get that full curve. But edge B, although it is above minVal and is in same region as that of edge C, it is not connected to any “sure-edge”, so that is discarded. So it is very important that we have to select minVal and maxVal accordingly to get the correct result.

This stage also removes small pixels noises on the assumption that edges are long lines.

So what we finally get is strong edges in the image.

double threshold to determine potential edges

Track edge by **hysteresis**: Finalize the detection of edges by suppressing all the other edges that are weak and not connected to strong edges.



Canny edge detector

- This is probably the most widely used edge detector in computer vision
- Theoretical model: step-edges corrupted by additive Gaussian noise
- Canny has shown that the first derivative of the Gaussian closely approximates the operator that optimizes the product of signal-to-noise ratio and localization

J. Canny, [A Computational Approach To Edge Detection](#), IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

Source: L. Fei-Fei



Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- Non-maximum suppression:
 - Thin multi-pixel wide "ridges" down to single pixel width
- Linking and thresholding (hysteresis):
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them

MATLAB: `edge(image, 'canny')`

Source: D. Lowe, L. Fei-Fei



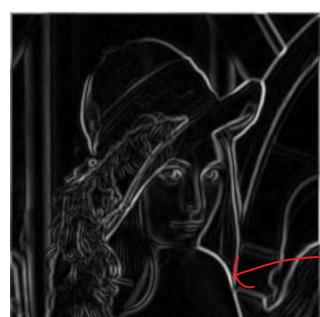
Example



original image (Lena)



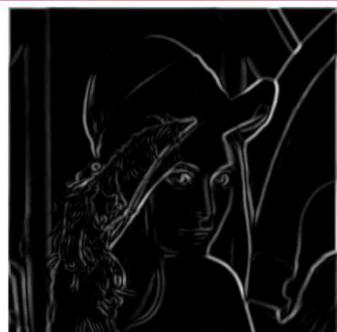
Example



norm of the gradient



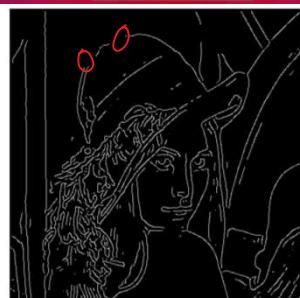
Example



thresholding



Example



thinning
(non-maximum suppression)

C_π Non-maximum suppression

Source: D. Forsyth

C_π Edge linking

Source: D. Forsyth

C_π Hysteresis thresholding

Check that if maximum value of gradient is sufficiently large

- drop-outs? use **hysteresis**
 - use a high threshold to start edge curves and a low threshold to continue them.

Source: S. Seitz

C_π Hysteresis thresholding

Source: L. Fei-Fei

Effect of σ (Gaussian kernel spread/size)

original Canny with $\sigma = 1$ Canny with $\sigma = 2$

The choice of σ depends on desired behavior

- large σ detects large scale edges
- small σ detects fine features

Source: S. Seitz

Edge detection is just the beginning...

Berkeley segmentation database:
<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

image	human segmentation	gradient magnitude